

SyriaTel Customer Churn Prediction and Retention Strategy

Author: John Mugambi

Overview

SyriaTel, a telecommunications company faces challenges in customer retention as some users discontinue their services. By leveraging machine learning techniques, we aim to predict customer churn and provide strategic insights to help SyriaTel improve customer retention.

Business Problem

Customer churn impacts SyriaTel revenue. The company needs an effective way to identify customers who are likely to leave and take proactive measures to retain them. Understanding the factors influencing churn will allow SyriaTel to enhance customer satisfaction, optimize marketing strategies, and improve overall business performance.

Objectives

Develop a machine learning model to predict whether a customer is likely to churn.

Identify key factors that influence customer churn.

Provide actionable insights to SyriaTel for improving customer retention strategies

Business Questions

1. What customer behaviors or attributes are most predictive of churn?
2. How can SyriaTel proactively intervene to retain at-risk customers?
3. Which classification models perform best in predicting customer churn?
4. What strategies can be implemented to enhance customer satisfaction and reduce churn?

Data Understanding

Data Sources and Relevance

The dataset for this project comes from Kaggle and contains customer data related to SyriaTel, a telecommunications company. This dataset is highly relevant to the problem of customer churn prediction. By analyzing customer behaviors and service usage patterns, we can identify key factors influencing churn. Understanding these factors will help SyriaTel develop effective strategies to retain customers, improve satisfaction, and reduce financial losses associated with customer attrition.

```
In [439]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from scipy.stats import boxcox
```

```
In [440]: # Load the data
churn_df = pd.read_csv("../data/churn_dataset.csv")
churn_df
```

Out[440]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...
...
3328	AZ	192	415	414-4276	no	yes	36	156.2	77	26.55	...
3329	WV	68	415	370-3271	no	no	0	231.1	57	39.29	...
3330	RI	28	510	328-8230	no	no	0	180.8	109	30.74	...
3331	CT	184	510	364-6381	yes	no	0	213.8	105	36.35	...
3332	TN	74	415	400-4344	no	yes	25	234.4	113	39.85	...

3333 rows × 21 columns



In [441]: churn_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   phone number                        3333 non-null   object
4   international plan                  3333 non-null   object
5   voice mail plan                     3333 non-null   object
6   number vmail messages               3333 non-null   int64
7   total day minutes                   3333 non-null   float64
8   total day calls                     3333 non-null   int64
9   total day charge                    3333 non-null   float64
10  total eve minutes                   3333 non-null   float64
11  total eve calls                     3333 non-null   int64
12  total eve charge                    3333 non-null   float64
13  total night minutes                 3333 non-null   float64
14  total night calls                   3333 non-null   int64
15  total night charge                  3333 non-null   float64
16  total intl minutes                  3333 non-null   float64
17  total intl calls                    3333 non-null   int64
18  total intl charge                   3333 non-null   float64
19  customer service calls              3333 non-null   int64
20  churn                              3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

In [442]: churn_df.shape

Out[442]: (3333, 21)

In [443]: churn_df.describe()

Out[443]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.98034
std	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.71384
min	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.600000
50%	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.400000
75%	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.300000
max	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.700000

```
In [444]: churn_df['state'].value_counts()
```

```
Out[444]: WV      106
          MN       84
          NY       83
          AL       80
          OR       78
          WI       78
          OH       78
          WY       77
          VA       77
          CT       74
          VT       73
          MI       73
          ID       73
          UT       72
          TX       72
          IN       71
          MD       70
          KS       70
          MT       68
          NC       68
          NJ       68
          NV       66
          CO       66
          WA       66
          RI       65
          MS       65
          MA       65
          AZ       64
          MO       63
          FL       63
          ME       62
          ND       62
          NM       62
          OK       61
          DE       61
          NE       61
          SD       60
          SC       60
          KY       59
          IL       58
          NH       56
          AR       55
          DC       54
          GA       54
          TN       53
          HI       53
          AK       52
          LA       51
          PA       45
          IA       44
          CA       34
          Name: state, dtype: int64
```

```
In [445]: churn_df['international plan'].value_counts()
```

```
Out[445]: no      3010  
         yes       323  
         Name: international plan, dtype: int64
```

```
In [446]: churn_df['voice mail plan'].value_counts()
```

```
Out[446]: no      2411  
         yes       922  
         Name: voice mail plan, dtype: int64
```

Data Preparation

Data Cleaning

```
In [447]: #Format the column names removing unconventional naming ways  
churn_df.columns = churn_df.columns.str.replace(' ', '_')  
churn_df.columns
```

```
Out[447]: Index(['state', 'account_length', 'area_code', 'phone_number',  
                'international_plan', 'voice_mail_plan', 'number_vmail_messages',  
                'total_day_minutes', 'total_day_calls', 'total_day_charge',  
                'total_eve_minutes', 'total_eve_calls', 'total_eve_charge',  
                'total_night_minutes', 'total_night_calls', 'total_night_charge',  
                'total_intl_minutes', 'total_intl_calls', 'total_intl_charge',  
                'customer_service_calls', 'churn'],  
              dtype='object')
```

```
In [448]: #Count number of missing values in the dataframe  
churn_df.isnull().sum()
```

```
Out[448]: state                                0  
account_length                             0  
area_code                                  0  
phone_number                              0  
international_plan                        0  
voice_mail_plan                          0  
number_vmail_messages                    0  
total_day_minutes                        0  
total_day_calls                         0  
total_day_charge                        0  
total_eve_minutes                       0  
total_eve_calls                        0  
total_eve_charge                        0  
total_night_minutes                    0  
total_night_calls                     0  
total_night_charge                    0  
total_intl_minutes                    0  
total_intl_calls                     0  
total_intl_charge                    0  
customer_service_calls                0  
churn                                  0  
dtype: int64
```

```
In [449]: # Check for duplicates  
churn_df.duplicated().sum()
```

```
Out[449]: 0
```

```
In [450]: # Converting the categorical columns to categorical types  
churn_df["international_plan"] = churn_df["international_plan"].astype("category")  
churn_df["voice_mail_plan"] = churn_df["voice_mail_plan"].astype("category")
```

```
In [451]: # Convert the boolean to int 0/1  
#churn_df["churn"] = churn_df["churn"].astype("int")
```

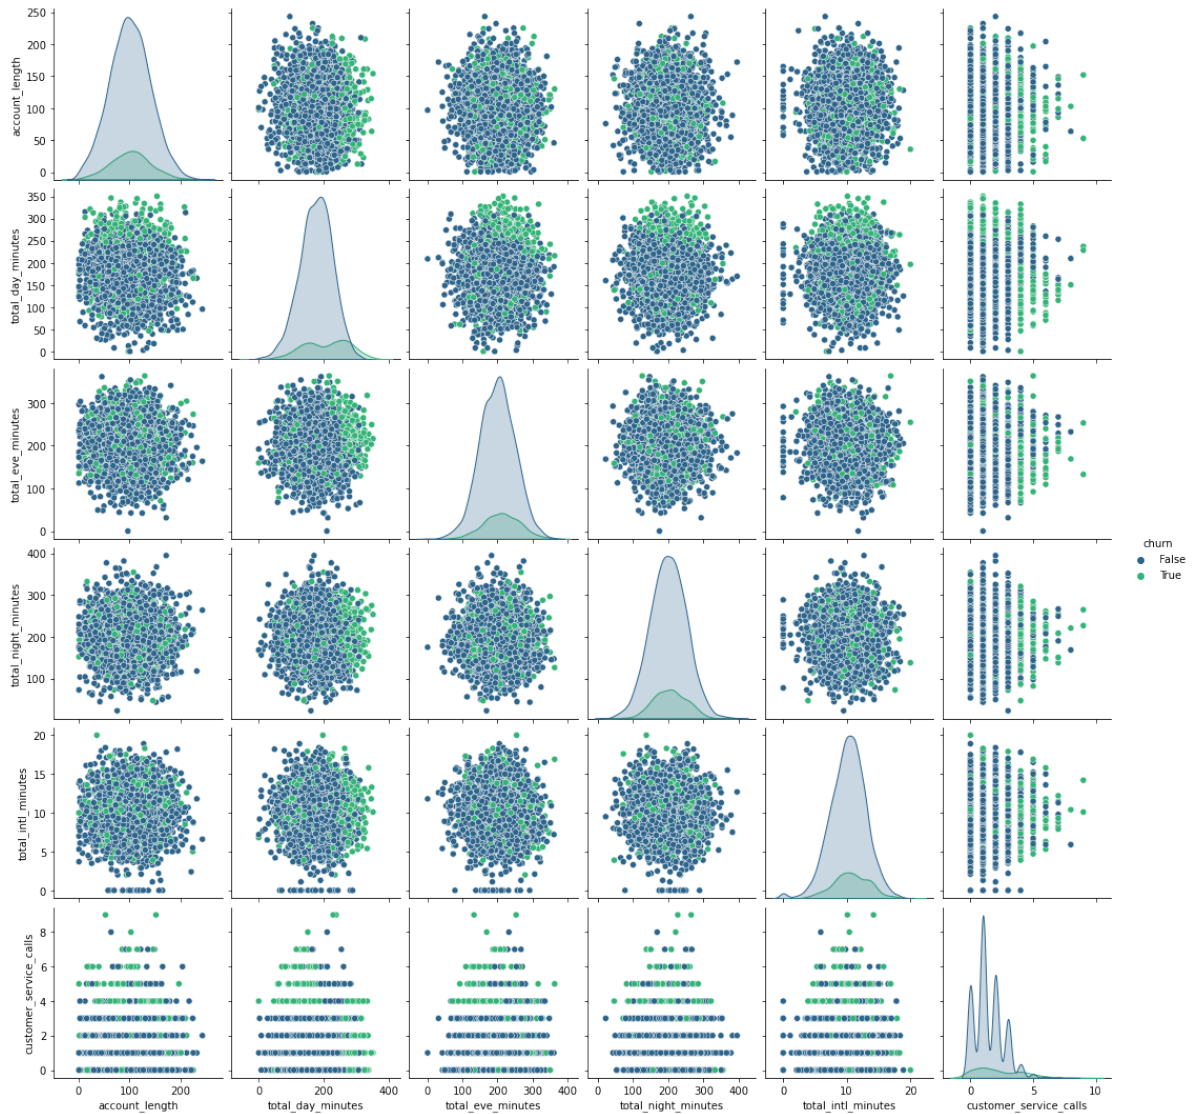
Feature Engineering

```
In [452]: #Drop the phone number column since it may not provide usefull predictive value  
churn_df.drop(columns=["phone_number"], inplace=True)
```

```
In [453]: #Create a new total calls feature  
churn_df["total_calls"] = (  
    churn_df["total_day_calls"]  
    + churn_df["total_eve_calls"]  
    + churn_df["total_night_calls"]  
    + churn_df["total_intl_calls"]  
)
```

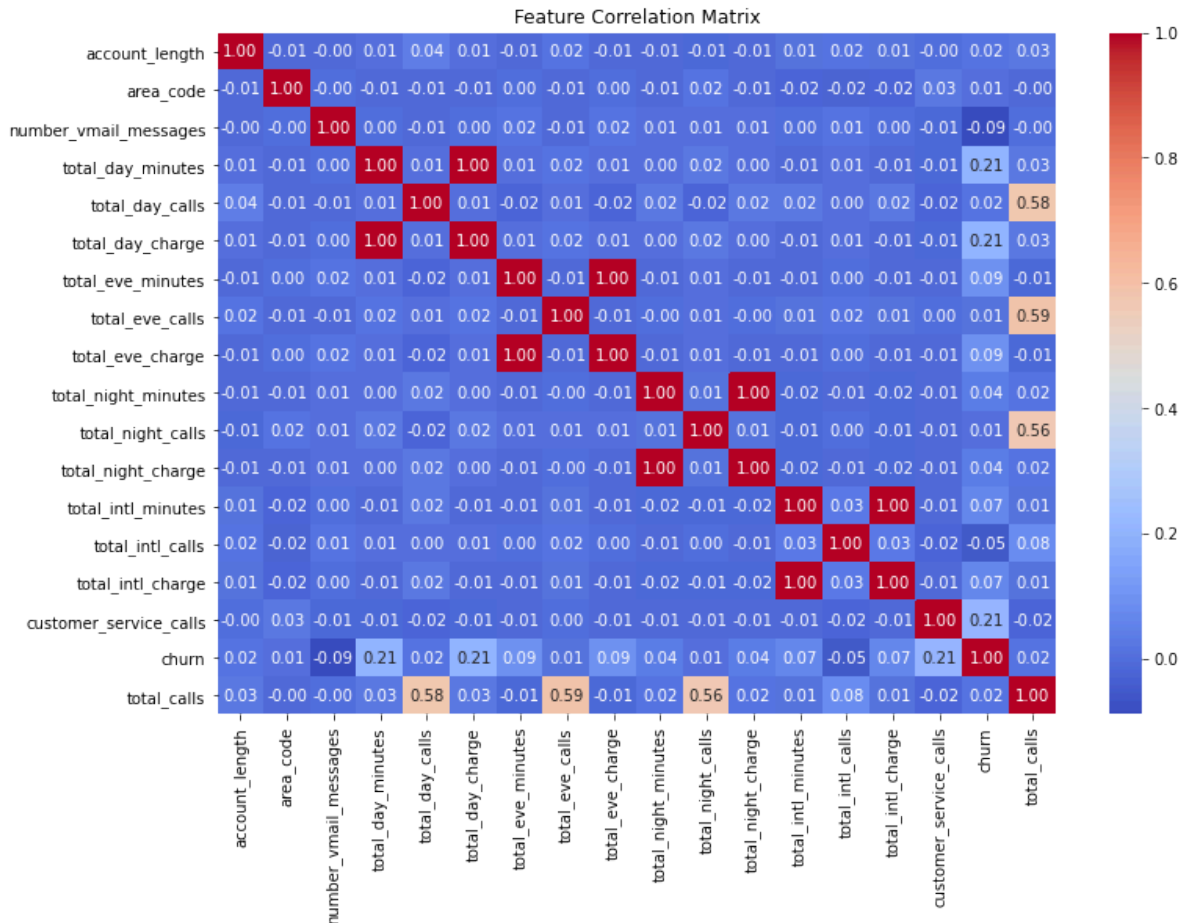
Visualization

```
In [454]: sns.pairplot(churn_df[['state', 'account_length', 'total_day_minutes', 'total_eve_minutes', 'total_night_minutes', 'total_intl_minutes', 'customer_service_calls'], hue='churn', palette='viridis', diag_kind='kde')
plt.show()
```




Data Modeling



```
In [455]: #Correlation Analysis
plt.figure(figsize=(12, 8))
sns.heatmap(churn_df.corr(), annot=True, fmt=".2f", cmap="coolwarm")
plt.title("Feature Correlation Matrix")
plt.show()
```




```
In [456]: #Check the churn distribution
sns.countplot(x=churn_df["churn"])
plt.title("Churn Distribution")
plt.show()
```




```
In [457]: #Boxplots for Numerical Variables:  
sns.boxplot(x="churn", y="total_day_minutes", data=churn_df)  
plt.show()
```




```
In [458]: #Visualize if customers service calls relate to churn  
plt.figure(figsize=(10, 6))  
sns.barplot(x='churn', y='customer_service_calls', data=churn_df, palette='viridis')  
plt.title('Average Customer Service Calls by Churn')  
plt.xlabel('Churn')  
plt.ylabel('Average Number of Customer Service Calls')  
plt.xticks([0, 1], ['Stayed', 'Churned'])  
plt.show()
```



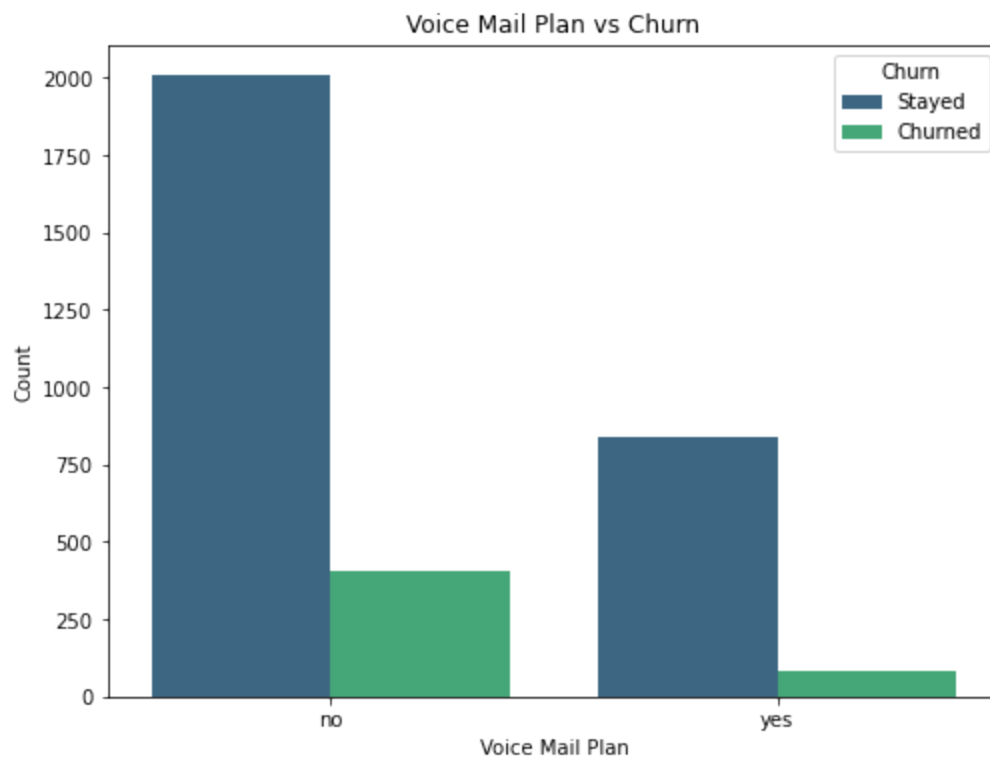
```
In [459]: plt.figure(figsize=(8, 6))
sns.countplot(x='international_plan', hue='churn', data=churn_df, palette='viridis')
plt.title('International Plan vs Churn')
plt.xlabel('International Plan')
plt.ylabel('Count')
plt.legend(title='Churn', labels=['Stayed', 'Churned'])
plt.show()
```



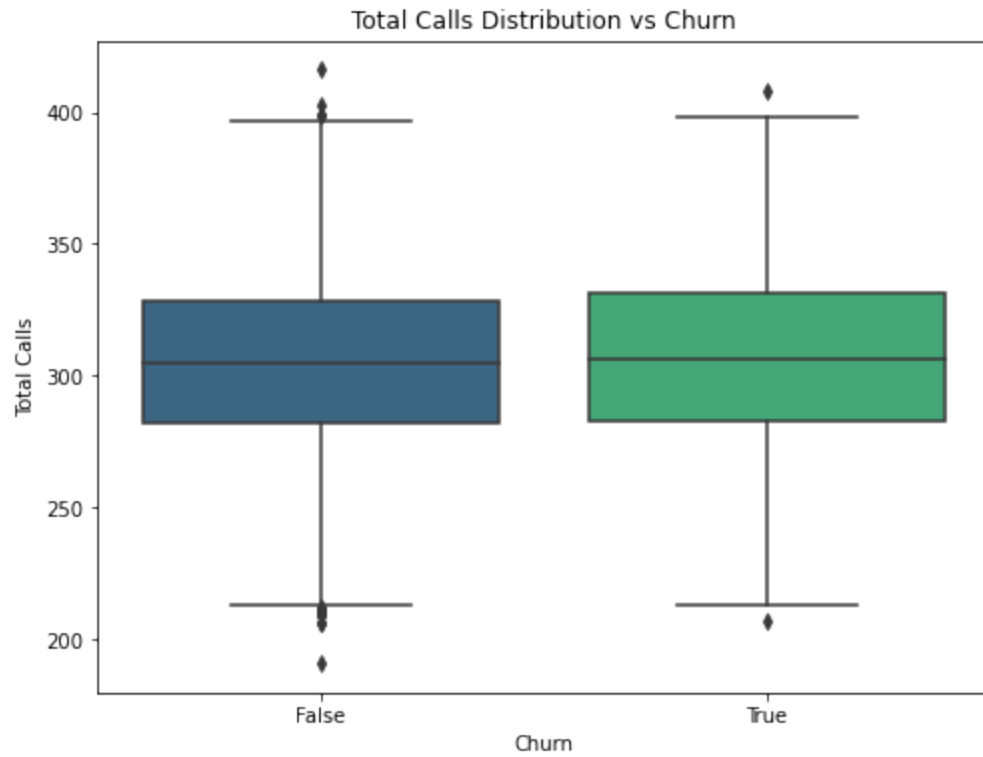
```
In [460]: plt.figure(figsize=(8, 6))
sns.countplot(x='area_code', hue='churn', data=churn_df, palette='viridis')
plt.title('Area Code vs Churn')
plt.xlabel('Area Code')
plt.ylabel('Count')
plt.legend(title='Churn', labels=['Stayed', 'Churned'])
plt.show()
```



```
In [461]: plt.figure(figsize=(8, 6))
sns.countplot(x='voice_mail_plan', hue='churn', data=churn_df, palette='viridis')
plt.title('Voice Mail Plan vs Churn')
plt.xlabel('Voice Mail Plan')
plt.ylabel('Count')
plt.legend(title='Churn', labels=['Stayed', 'Churned'])
plt.show()
```



```
In [462]: #total calls distribution
plt.figure(figsize=(8, 6))
sns.boxplot(x='churn', y='total_calls', data=churn_df, palette='viridis')
plt.title('Total Calls Distribution vs Churn')
plt.xlabel('Churn')
plt.ylabel('Total Calls')
plt.show()
```



Preprocessing

In [463]: churn_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account_length                       3333 non-null   int64
2   area_code                            3333 non-null   int64
3   international_plan                   3333 non-null   category
4   voice_mail_plan                      3333 non-null   category
5   number_vmail_messages                3333 non-null   int64
6   total_day_minutes                    3333 non-null   float64
7   total_day_calls                      3333 non-null   int64
8   total_day_charge                     3333 non-null   float64
9   total_eve_minutes                    3333 non-null   float64
10  total_eve_calls                      3333 non-null   int64
11  total_eve_charge                     3333 non-null   float64
12  total_night_minutes                  3333 non-null   float64
13  total_night_calls                    3333 non-null   int64
14  total_night_charge                   3333 non-null   float64
15  total_intl_minutes                   3333 non-null   float64
16  total_intl_calls                     3333 non-null   int64
17  total_intl_charge                     3333 non-null   float64
18  customer_service_calls               3333 non-null   int64
19  churn                                3333 non-null   bool
20  total_calls                          3333 non-null   int64
dtypes: bool(1), category(2), float64(8), int64(9), object(1)
memory usage: 478.8+ KB
```

```
In [464]: # get the numeric cols and the cat cols
num_original_columns = churn_df.select_dtypes(include=np.number).columns.tolist()
categorical_cols = churn_df.select_dtypes(exclude=np.number).columns.tolist()

print("Categorical columns:", categorical_cols)
num_original_columns
```

Categorical columns: ['state', 'international_plan', 'voice_mail_plan', 'churn']

```
Out[464]: ['account_length',
'area_code',
'number_vmail_messages',
'total_day_minutes',
'total_day_calls',
'total_day_charge',
'total_eve_minutes',
'total_eve_calls',
'total_eve_charge',
'total_night_minutes',
'total_night_calls',
'total_night_charge',
'total_intl_minutes',
'total_intl_calls',
'total_intl_charge',
'customer_service_calls',
'total_calls']
```

Encoding

```
In [465]: churn_df[categorical_cols].info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   state                 3333 non-null   object
1   international_plan     3333 non-null   category
2   voice_mail_plan       3333 non-null   category
3   churn                 3333 non-null   bool
dtypes: bool(1), category(2), object(1)
memory usage: 36.1+ KB
```



```
In [466]: for i in categorical_cols:
           print(f'The variable "{i}" has {churn_df[i].nunique()} variables: {churn_df[i].unique()})
```

```
The variable "state" has 51 variables: ['KS' 'OH' 'NJ' 'OK' 'AL' 'MA' 'MO' 'LA' 'WV' 'IN' 'RI' 'IA' 'MT' 'NY'
'ID' 'VT' 'VA' 'TX' 'FL' 'CO' 'AZ' 'SC' 'NE' 'WY' 'HI' 'IL' 'NH' 'GA'
'AK' 'MD' 'AR' 'WI' 'OR' 'MI' 'DE' 'UT' 'CA' 'MN' 'SD' 'NC' 'WA' 'NM'
'NV' 'DC' 'KY' 'ME' 'MS' 'TN' 'PA' 'CT' 'ND']
```

```
The variable "international_plan" has 2 variables: ['no', 'yes']
Categories (2, object): ['no', 'yes']
```

```
The variable "voice_mail_plan" has 2 variables: ['yes', 'no']
Categories (2, object): ['yes', 'no']
```

```
The variable "churn" has 2 variables: [False True]
```

```
In [467]: # Initialize Label Encoder
le = LabelEncoder()
```

```
In [468]: # Apply Label Encoding to binary categorical columns
churn_df["international_plan"] = le.fit_transform(churn_df["international_plan"])
churn_df["voice_mail_plan"] = le.fit_transform(churn_df["voice_mail_plan"])
```

```
In [469]: # One-Hot Encode 'state' column, and drop it from the dataframe
churn_df = pd.get_dummies(churn_df, columns=["state"], drop_first=True)
```

In [470]: churn_df.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3333 entries, 0 to 3332
```

```
Data columns (total 70 columns):
```

#	Column	Non-Null Count	Dtype
0	account_length	3333 non-null	int64
1	area_code	3333 non-null	int64
2	international_plan	3333 non-null	int32
3	voice_mail_plan	3333 non-null	int32
4	number_vmail_messages	3333 non-null	int64
5	total_day_minutes	3333 non-null	float64
6	total_day_calls	3333 non-null	int64
7	total_day_charge	3333 non-null	float64
8	total_eve_minutes	3333 non-null	float64
9	total_eve_calls	3333 non-null	int64
10	total_eve_charge	3333 non-null	float64
11	total_night_minutes	3333 non-null	float64
12	total_night_calls	3333 non-null	int64
13	total_night_charge	3333 non-null	float64
14	total_intl_minutes	3333 non-null	float64
15	total_intl_calls	3333 non-null	int64
16	total_intl_charge	3333 non-null	float64
17	customer_service_calls	3333 non-null	int64
18	churn	3333 non-null	bool
19	total_calls	3333 non-null	int64
20	state_AL	3333 non-null	uint8
21	state_AR	3333 non-null	uint8
22	state_AZ	3333 non-null	uint8
23	state_CA	3333 non-null	uint8
24	state_CO	3333 non-null	uint8
25	state_CT	3333 non-null	uint8
26	state_DC	3333 non-null	uint8
27	state_DE	3333 non-null	uint8
28	state_FL	3333 non-null	uint8
29	state_GA	3333 non-null	uint8
30	state_HI	3333 non-null	uint8
31	state_IA	3333 non-null	uint8
32	state_ID	3333 non-null	uint8
33	state_IL	3333 non-null	uint8
34	state_IN	3333 non-null	uint8
35	state_KS	3333 non-null	uint8
36	state_KY	3333 non-null	uint8
37	state_LA	3333 non-null	uint8
38	state_MA	3333 non-null	uint8
39	state_MD	3333 non-null	uint8
40	state_ME	3333 non-null	uint8
41	state_MI	3333 non-null	uint8
42	state_MN	3333 non-null	uint8
43	state_MO	3333 non-null	uint8
44	state_MS	3333 non-null	uint8
45	state_MT	3333 non-null	uint8
46	state_NC	3333 non-null	uint8
47	state_ND	3333 non-null	uint8
48	state_NE	3333 non-null	uint8
49	state_NH	3333 non-null	uint8
50	state_NJ	3333 non-null	uint8
51	state_NM	3333 non-null	uint8

```

52 state_NV      3333 non-null  uint8
53 state_NY      3333 non-null  uint8
54 state_OH      3333 non-null  uint8
55 state_OK      3333 non-null  uint8
56 state_OR      3333 non-null  uint8
57 state_PA      3333 non-null  uint8
58 state_RI      3333 non-null  uint8
59 state_SC      3333 non-null  uint8
60 state_SD      3333 non-null  uint8
61 state_TN      3333 non-null  uint8
62 state_TX      3333 non-null  uint8
63 state_UT      3333 non-null  uint8
64 state_VA      3333 non-null  uint8
65 state_VT      3333 non-null  uint8
66 state_WA      3333 non-null  uint8
67 state_WI      3333 non-null  uint8
68 state_WV      3333 non-null  uint8
69 state_WY      3333 non-null  uint8
dtypes: bool(1), float64(8), int32(2), int64(9), uint8(50)
memory usage: 634.8 KB


```

In [471]: churn_df.head()

Out[471]:

	account_length	area_code	international_plan	voice_mail_plan	number_vmail_messages	total
0	128	415	0	1		25
1	107	415	0	1		26
2	137	415	0	0		0
3	84	408	1	0		0
4	75	415	1	0		0

5 rows × 70 columns



Scaling

In [472]: scaler = StandardScaler()

In [473]: *# Select numerical columns without the encoded categorical ones*
numerical_cols = churn_df.select_dtypes(include=np.number).columns

```
In [474]: #Check Skewness
for col in num_original_columns:
    skew_value = churn_df[col].skew()
    print(f"Column: {col}, Skewness: {skew_value:.4f}")
```

```
Column: account_length, Skewness: 0.0966
Column: area_code, Skewness: 1.1268
Column: number_vmail_messages, Skewness: 1.2648
Column: total_day_minutes, Skewness: -0.0291
Column: total_day_calls, Skewness: -0.1118
Column: total_day_charge, Skewness: -0.0291
Column: total_eve_minutes, Skewness: -0.0239
Column: total_eve_calls, Skewness: -0.0556
Column: total_eve_charge, Skewness: -0.0239
Column: total_night_minutes, Skewness: 0.0089
Column: total_night_calls, Skewness: 0.0325
Column: total_night_charge, Skewness: 0.0089
Column: total_intl_minutes, Skewness: -0.2451
Column: total_intl_calls, Skewness: 1.3215
Column: total_intl_charge, Skewness: -0.2453
Column: customer_service_calls, Skewness: 1.0914
Column: total_calls, Skewness: -0.0376
```

Based on the above the following are Highly skewed (consider transformation):

- area_code (1.1268)
- number_vmail_messages (1.2648)
- total_intl_calls (1.3215)
- customer_service_calls (1.0914)
- churn (2.0184)

Moderately skewed:

- total_intl_minutes (-0.2451)
- total_intl_charge (-0.2453)

```
In [475]: # Box-Cox transformation was chosen for right-skewed
# data like number_vmail_messages, total_intl_calls, customer_service_calls
# i.e values greater than 1

for col in ["number_vmail_messages", "total_intl_calls", "customer_service_calls"]:
    churn_df[col] = boxcox(churn_df[col] + 1)
```

```
In [476]: # Log transform 'churn' - has binary values
# churn_df["churn"] = np.log1p(churn_df["churn"])
```

```
In [477]: #Confirm skewness has improved
for col in ["number_vmail_messages", "total_intl_calls", "customer_service_calls"]:
    print(f"Column: {col}, Skewness: {churn_df[col].skew()}")
```


```
Column: number_vmail_messages, Skewness: 1.0002251990902429
Column: total_intl_calls, Skewness: 0.005816369249351684
Column: customer_service_calls, Skewness: -0.013769770295669267
```

```
In [478]: # Apply scaling
churn_df_scaled = churn_df.copy()
churn_df_scaled[numerical_cols] = scaler.fit_transform(churn_df[numerical_cols])
churn_df_scaled.head()
```

Out[478]:

	account_length	area_code	international_plan	voice_mail_plan	number_vmail_messages	total
0	0.676489	-0.523603	-0.327580	1.617086		1.608479
1	0.149065	-0.523603	-0.327580	1.617086		1.612570
2	0.902529	-0.523603	-0.327580	-0.618396		-0.618292
3	-0.428590	-0.688834	3.052685	-0.618396		-0.618292
4	-0.654629	-0.523603	3.052685	-0.618396		-0.618292

5 rows × 70 columns



```
In [479]: churn_df[numerical_cols].describe()
churn_df_scaled['churn']

numerical_cols
```

```
Out[479]: Index(['account_length', 'area_code', 'international_plan', 'voice_mail_pla
n',
               'number_vmail_messages', 'total_day_minutes', 'total_day_calls',
               'total_day_charge', 'total_eve_minutes', 'total_eve_calls',
               'total_eve_charge', 'total_night_minutes', 'total_night_calls',
               'total_night_charge', 'total_intl_minutes', 'total_intl_calls',
               'total_intl_charge', 'customer_service_calls', 'total_calls',
               'state_AL', 'state_AR', 'state_AZ', 'state_CA', 'state_CO', 'state_C
T',
               'state_DC', 'state_DE', 'state_FL', 'state_GA', 'state_HI', 'state_I
A',
               'state_ID', 'state_IL', 'state_IN', 'state_KS', 'state_KY', 'state_L
A',
               'state_MA', 'state_MD', 'state_ME', 'state_MI', 'state_MN', 'state_M
O',
               'state_MS', 'state_MT', 'state_NC', 'state_ND', 'state_NE', 'state_N
H',
               'state_NJ', 'state_NM', 'state_NV', 'state_NY', 'state_OH', 'state_O
K',
               'state_OR', 'state_PA', 'state_RI', 'state_SC', 'state_SD', 'state_T
N',
               'state_TX', 'state_UT', 'state_VA', 'state_VT', 'state_WA', 'state_W
I',
               'state_WV', 'state_WY'],
              dtype='object')
```

Modeling

```
In [480]: from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
import statsmodels.api as sm

from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error,
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [481]: # Defining dependent and independent variable  
X = churn_df[numerical_cols]  
y = churn_df['churn']  
  
X = sm.add_constant(X)  
model = sm.OLS(y, X).fit()  
  
print(model.summary())
```


OLS Regression Results

```

=====
=
Dep. Variable:          churn    R-squared:                0.18
5
Model:                  OLS      Adj. R-squared:           0.16
8
Method:                 Least Squares    F-statistic:             10.7
2
Date:                   Sat, 08 Mar 2025    Prob (F-statistic):      6.20e-10
0
Time:                   20:25:06    Log-Likelihood:          -908.9
4
No. Observations:      3333    AIC:                     195
8.
Df Residuals:          3263    BIC:                     238
6.
Df Model:               69
Covariance Type:       nonrobust
=====
=====

```

```

=====
=====
                                coef    std err          t      P>|t|      [0.02
5      0.975]
-----
-----
const                -0.3711      0.103     -3.610      0.000     -0.57
3      -0.170
account_length        0.0001      0.000      1.055      0.291     -0.00
0      0.000
area_code            -8.911e-05      0.000     -0.671      0.502     -0.00
0      0.000
international_plan    0.3056      0.019     16.008      0.000      0.26
8      0.343
voice_mail_plan      -1.0226      0.685     -1.494      0.135     -2.36
5      0.320
number_vmail_messages 0.9003      0.654      1.376      0.169     -0.38
2      2.183
total_day_minutes     0.0369      0.333      0.111      0.912     -0.61
7      0.691
total_day_calls       -0.0216      0.008     -2.560      0.010     -0.03
8      -0.005
total_day_charge      -0.2100      1.962     -0.107      0.915     -4.05
6      3.636
total_eve_minutes     0.0842      0.166      0.509      0.611     -0.24
0      0.409
total_eve_calls       -0.0219      0.008     -2.591      0.010     -0.03
8      -0.005
total_eve_charge      -0.9828      1.947     -0.505      0.614     -4.80
1      2.835
total_night_minutes   -0.0544      0.088     -0.615      0.539     -0.22
8      0.119
total_night_calls     -0.0220      0.008     -2.604      0.009     -0.03
9      -0.005
total_night_charge     1.2162      1.966      0.619      0.536     -2.63
9      5.072
total_intl_minutes    -0.4130      0.529     -0.780      0.435     -1.45
1      0.625

```

total_intl_calls	-0.1425	0.038	-3.731	0.000	-0.21
7 -0.068					
total_intl_charge	1.5610	1.960	0.796	0.426	-2.28
2 5.404					
customer_service_calls	0.0993	0.010	10.120	0.000	0.08
0 0.119					
total_calls	0.0220	0.008	2.604	0.009	0.00
5 0.038					
state_AL	0.0146	0.057	0.255	0.798	-0.09
8 0.127					
state_AR	0.0837	0.062	1.344	0.179	-0.03
8 0.206					
state_AZ	0.0055	0.060	0.091	0.927	-0.11
2 0.123					
state_CA	0.1738	0.071	2.447	0.014	0.03
5 0.313					
state_CO	0.0473	0.060	0.793	0.428	-0.07
0 0.164					
state_CT	0.0711	0.058	1.220	0.223	-0.04
3 0.185					
state_DC	0.0224	0.063	0.358	0.720	-0.10
0 0.145					
state_DE	0.0324	0.061	0.532	0.594	-0.08
7 0.152					
state_FL	0.0260	0.060	0.430	0.667	-0.09
3 0.145					
state_GA	0.0502	0.063	0.802	0.423	-0.07
3 0.173					
state_HI	-0.0091	0.063	-0.144	0.885	-0.13
2 0.114					
state_IA	0.0204	0.066	0.309	0.757	-0.10
9 0.150					
state_ID	0.0505	0.058	0.865	0.387	-0.06
4 0.165					
state_IL	-0.0245	0.062	-0.397	0.691	-0.14
5 0.096					
state_IN	0.0208	0.059	0.354	0.723	-0.09
5 0.136					
state_KS	0.0707	0.059	1.199	0.230	-0.04
5 0.186					
state_KY	0.0602	0.061	0.984	0.325	-0.06
0 0.180					
state_LA	0.0230	0.063	0.362	0.717	-0.10
1 0.147					
state_MA	0.0832	0.060	1.388	0.165	-0.03
4 0.201					
state_MD	0.1111	0.059	1.883	0.060	-0.00
5 0.227					
state_ME	0.1113	0.061	1.837	0.066	-0.00
7 0.230					
state_MI	0.1156	0.058	1.977	0.048	0.00
1 0.230					
state_MN	0.0856	0.057	1.506	0.132	-0.02
6 0.197					
state_MO	0.0332	0.060	0.551	0.582	-0.08
5 0.152					
state_MS	0.1163	0.060	1.942	0.052	-0.00

1	0.234					
state_MT		0.1522	0.059	2.568	0.010	0.03
6	0.268					
state_NC		0.0446	0.059	0.751	0.453	-0.07
2	0.161					
state_ND		-0.0138	0.061	-0.227	0.820	-0.13
3	0.105					
state_NE		0.0099	0.061	0.162	0.871	-0.10
9	0.129					
state_NH		0.0884	0.062	1.427	0.154	-0.03
3	0.210					
state_NJ		0.1520	0.059	2.561	0.010	0.03
6	0.268					
state_NM		-0.0017	0.061	-0.027	0.978	-0.12
1	0.117					
state_NV		0.0963	0.060	1.613	0.107	-0.02
1	0.213					
state_NY		0.0890	0.057	1.564	0.118	-0.02
3	0.201					
state_OH		0.0393	0.058	0.682	0.495	-0.07
4	0.152					
state_OK		0.0542	0.061	0.891	0.373	-0.06
5	0.174					
state_OR		0.0474	0.058	0.823	0.411	-0.06
6	0.160					
state_PA		0.0788	0.066	1.202	0.229	-0.05
0	0.207					
state_RI		0.0095	0.060	0.159	0.874	-0.10
8	0.127					
state_SC		0.1841	0.061	3.015	0.003	0.06
4	0.304					
state_SD		0.0445	0.061	0.730	0.465	-0.07
5	0.164					
state_TN		0.0095	0.063	0.151	0.880	-0.11
4	0.133					
state_TX		0.1556	0.059	2.657	0.008	0.04
1	0.270					
state_UT		0.0703	0.059	1.202	0.229	-0.04
4	0.185					
state_VA		-0.0429	0.058	-0.742	0.458	-0.15
6	0.071					
state_VT		0.0078	0.058	0.134	0.894	-0.10
7	0.122					
state_WA		0.1316	0.060	2.204	0.028	0.01
5	0.249					
state_WI		0.0129	0.058	0.223	0.823	-0.10
0	0.126					
state_WV		0.0311	0.055	0.571	0.568	-0.07
6	0.138					
state_WY		0.0165	0.058	0.285	0.775	-0.09
7	0.130					

=====

Omnibus:	844.862	Durbin-Watson:	1.96
8			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1690.78
8			

Skew:	1.531	Prob(JB):	0.0
0			
Kurtosis:	4.673	Cond. No.	2.43e+0
5			

=====

=

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.43e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [482]: # Make predictions
y_pred = model.predict(X)

# Calculate the metrics
rmse = np.sqrt(mean_squared_error(y, y_pred))
r2 = r2_score(y, y_pred)
mae = mean_absolute_error(y, y_pred)

print(f"\nRMSE: {rmse}")
print(f"R-squared: {r2}")
print(f"Mean Absolute Error: {mae}")
```

```
RMSE: 0.3178328250152115
R-squared: 0.1847775467138406
Mean Absolute Error: 0.22132800557877938
```

While the model seems to fit perfectly ($R^2 = 1.0$, RMSE close to zero)

The F-statistic is 7.63e+27, which is extremely high, and the F-statistic= 0.00, indicates that the overall model is statistically significant.

The RMSE value is very low suggesting that the model's predictions are very close to the actual values.

The skewness value is -0.042, which is very close to 0, therefore meaning the model doesn't have large biases.

Classification

Logistic Regression

```
In [483]: from sklearn.model_selection import train_test_split

# Define features (X) and target (y)
X = churn_df_scaled.drop(columns=['churn'])
y = churn_df_scaled['churn']

# Split data into training and testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random
```

```
In [484]: print(y.unique())
churn_df_scaled['churn']
```

```
[False  True]
```

```
Out[484]: 0      False
1      False
2      False
3      False
4      False
...
3328   False
3329   False
3330   False
3331   False
3332   False
Name: churn, Length: 3333, dtype: bool
```

```

In [485]: # Initialize, train logistic regression model
logreg_model = LogisticRegression(max_iter=1000)
logreg_model.fit(X_train, y_train)

y_pred = logreg_model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print('Model 1 - Logistic Regression')
print(f"Accuracy: {accuracy}")
print(f"Recall: {recall}")
print(f"Precision: {precision}")
print(f"F1-score: {f1}")

# Classification report
print(classification_report(y_test, y_pred))

# Plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted No Churn (False)', 'Predicted Churn (True)',
            yticklabels=['Actual No Churn (False)', 'Actual Churn (True)'])
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

Model 1 - Logistic Regression


Accuracy: 0.8545727136431784

Recall: 0.18811881188118812

Precision: 0.5588235294117647

F1-score: 0.2814814814814815

	precision	recall	f1-score	support
False	0.87	0.97	0.92	566
True	0.56	0.19	0.28	101
accuracy			0.85	667
macro avg	0.71	0.58	0.60	667
weighted avg	0.82	0.85	0.82	667



Recall is quite low at 0.19, indicating that many customers who actually churned were not identified by the model. We will try other classification algorithms that might handle the imbalance better

```

In [486]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
import seaborn as sns
import matplotlib.pyplot as plt

# Initialize, train Decision Tree model
dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred_dt = dt_model.predict(X_test)

# Evaluate the model
accuracy_dt = accuracy_score(y_test, y_pred_dt)
recall_dt = recall_score(y_test, y_pred_dt)
precision_dt = precision_score(y_test, y_pred_dt)
f1_dt = f1_score(y_test, y_pred_dt)

print('Model 2 - Decision Tree Classifier')
print(f"Accuracy: {accuracy_dt}")
print(f"Recall: {recall_dt}")
print(f"Precision: {precision_dt}")
print(f"F1-score: {f1_dt}")

# Classification report
print(classification_report(y_test, y_pred_dt))

# Plot confusion matrix
cm_dt = confusion_matrix(y_test, y_pred_dt)
plt.figure(figsize=(6,5))
sns.heatmap(cm_dt, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted No Churn (False)', 'Predicted Churn (True)',
            yticklabels=['Actual No Churn (False)', 'Actual Churn (True)'])
plt.title("Confusion Matrix - Decision Tree")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

Model 2 - Decision Tree Classifier


Accuracy: 0.9235382308845578

Recall: 0.7425742574257426

Precision: 0.75

F1-score: 0.746268656716418

	precision	recall	f1-score	support
False	0.95	0.96	0.95	566
True	0.75	0.74	0.75	101
accuracy			0.92	667
macro avg	0.85	0.85	0.85	667
weighted avg	0.92	0.92	0.92	667



```

In [487]: from sklearn.ensemble import RandomForestClassifier

# Initialize, train Random Forest model
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)

y_pred_rf = rf_model.predict(X_test)

# Evaluate the model
accuracy_rf = accuracy_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf)

print('Model 3 - Random Forest Classifier')
print(f"Accuracy: {accuracy_rf}")
print(f"Recall: {recall_rf}")
print(f"Precision: {precision_rf}")
print(f"F1-score: {f1_rf}")

# Classification report
print(classification_report(y_test, y_pred_rf))

# Plot confusion matrix
cm_rf = confusion_matrix(y_test, y_pred_rf)
plt.figure(figsize=(6,5))
sns.heatmap(cm_rf, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted No Churn (False)', 'Predicted Churn (True)',
            yticklabels=['Actual No Churn (False)', 'Actual Churn (True)'])
plt.title("Confusion Matrix - Random Forest")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

Model 3 - Random Forest Classifier


Accuracy: 0.9385307346326837

Recall: 0.6138613861386139

Precision: 0.96875

F1-score: 0.7515151515151515

	precision	recall	f1-score	support
False	0.94	1.00	0.96	566
True	0.97	0.61	0.75	101
accuracy			0.94	667
macro avg	0.95	0.81	0.86	667
weighted avg	0.94	0.94	0.93	667



```
In [488]: from sklearn.neighbors import KNeighborsClassifier

# Initialize, train KNN model
knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train, y_train)

y_pred_knn = knn_model.predict(X_test)

# Evaluate the model
accuracy_knn = accuracy_score(y_test, y_pred_knn)
recall_knn = recall_score(y_test, y_pred_knn)
precision_knn = precision_score(y_test, y_pred_knn)
f1_knn = f1_score(y_test, y_pred_knn)

print('Model 4 - KNN Classifier')
print(f"Accuracy: {accuracy_knn}")
print(f"Recall: {recall_knn}")
print(f"Precision: {precision_knn}")
print(f"F1-score: {f1_knn}")

# Classification report
print(classification_report(y_test, y_pred_knn))

# Plot confusion matrix
cm_knn = confusion_matrix(y_test, y_pred_knn)
plt.figure(figsize=(6,5))
sns.heatmap(cm_knn, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted No Churn (False)', 'Predicted Churn (True)',
            yticklabels=['Actual No Churn (False)', 'Actual Churn (True)'])
plt.title("Confusion Matrix - KNN")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

Model 4 - KNN Classifier


Accuracy: 0.848575712143928

Recall: 0.04950495049504951

Precision: 0.5

F1-score: 0.09009009009009009

	precision	recall	f1-score	support
False	0.85	0.99	0.92	566
True	0.50	0.05	0.09	101
accuracy			0.85	667
macro avg	0.68	0.52	0.50	667
weighted avg	0.80	0.85	0.79	667



```
In [489]: from sklearn.svm import SVC

# Initialize, train SVM model
svm_model = SVC(kernel='linear', random_state=42) # You can also experiment w
svm_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred_svm = svm_model.predict(X_test)

# Evaluate the model
accuracy_svm = accuracy_score(y_test, y_pred_svm)
recall_svm = recall_score(y_test, y_pred_svm)
precision_svm = precision_score(y_test, y_pred_svm)
f1_svm = f1_score(y_test, y_pred_svm)

print('Model 5 - SVM Classifier')
print(f"Accuracy: {accuracy_svm}")
print(f"Recall: {recall_svm}")
print(f"Precision: {precision_svm}")
print(f"F1-score: {f1_svm}")

# Classification report
print(classification_report(y_test, y_pred_svm))

# Plot confusion matrix
cm_svm = confusion_matrix(y_test, y_pred_svm)
plt.figure(figsize=(6,5))
sns.heatmap(cm_svm, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted No Churn (False)', 'Predicted Churn (True)',
            yticklabels=['Actual No Churn (False)', 'Actual Churn (True)'])
plt.title("Confusion Matrix - SVM")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

C:\Users\jmuriithi\.conda\envs\learn-env\lib\site-packages\sklearn\metrics_classification.py:1221: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.


_warn_prf(average, modifier, msg_start, len(result))

C:\Users\jmuriithi\.conda\envs\learn-env\lib\site-packages\sklearn\metrics_classification.py:1221: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

Model 5 - SVM Classifier
Accuracy: 0.848575712143928
Recall: 0.0
Precision: 0.0
F1-score: 0.0

	precision	recall	f1-score	support
False	0.85	1.00	0.92	566
True	0.00	0.00	0.00	101
accuracy			0.85	667
macro avg	0.42	0.50	0.46	667
weighted avg	0.72	0.85	0.78	667



```
In [490]: from sklearn.naive_bayes import GaussianNB

# Initialize, train Naive Bayes model
nb_model = GaussianNB()
nb_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred_nb = nb_model.predict(X_test)

# Evaluate the model
accuracy_nb = accuracy_score(y_test, y_pred_nb)
recall_nb = recall_score(y_test, y_pred_nb)
precision_nb = precision_score(y_test, y_pred_nb)
f1_nb = f1_score(y_test, y_pred_nb)

print('Model 6 - Naive Bayes Classifier')
print(f"Accuracy: {accuracy_nb}")
print(f"Recall: {recall_nb}")
print(f"Precision: {precision_nb}")
print(f"F1-score: {f1_nb}")
print(classification_report(y_test, y_pred_nb))

# Plot confusion matrix
cm_nb = confusion_matrix(y_test, y_pred_nb)
plt.figure(figsize=(6,5))
sns.heatmap(cm_nb, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted No Churn (False)', 'Predicted Churn (True)',
            yticklabels=['Actual No Churn (False)', 'Actual Churn (True)'])
plt.title("Confusion Matrix - Naive Bayes")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

Model 6 - Naive Bayes Classifier


Accuracy: 0.568215892053973

Recall: 0.5643564356435643

Precision: 0.1893687707641196

F1-score: 0.2835820895522388

	precision	recall	f1-score	support
False	0.88	0.57	0.69	566
True	0.19	0.56	0.28	101
accuracy			0.57	667
macro avg	0.53	0.57	0.49	667
weighted avg	0.78	0.57	0.63	667



Best Performing Model

- Random Forest shows the best balance of accuracy (93.85%), precision (96.88%), and F1-score (75.15%), despite having a slightly lower recall (61.39%) compared to the decision tree.

Second Best Model

- Decision Tree also performs well with high accuracy (92.35%), decent recall (74.26%), and a good F1-score (74.63%). It's a great model for detecting churn.

Weak Models

- KNN and SVM fail to predict churn effectively, particularly in terms of recall and precision.

Hyperparameter Tuning

```
In [491]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
import matplotlib.pyplot as plt
import seaborn as sns

# Parameter grid for Random Forest
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Initialize the RandomForest model
rf_model = RandomForestClassifier(random_state=42)
# Initialize GridSearchCV
rf_grid = GridSearchCV(estimator=rf_model, param_grid=param_grid, cv=5, verbose=1)

rf_grid.fit(X_train, y_train)
rf_model_tuned = rf_grid.best_estimator_
y_pred = rf_model_tuned.predict(X_test)

# Evaluate the tuned model
rf_accuracy = accuracy_score(y_test, y_pred)
rf_recall = recall_score(y_test, y_pred)
rf_precision = precision_score(y_test, y_pred)
rf_f1 = f1_score(y_test, y_pred)

print(f"Tuned Random Forest - Accuracy: {rf_accuracy}")
print(f"Tuned Random Forest - Recall: {rf_recall}")
print(f"Tuned Random Forest - Precision: {rf_precision}")
print(f"Tuned Random Forest - F1-score: {rf_f1}")
print(classification_report(y_test, y_pred))


# Confusion matrix for the tuned model
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted False', 'Predicted True'],
            yticklabels=['Actual False', 'Actual True'])
plt.title("Confusion Matrix (Tuned Random Forest)")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

Fitting 5 folds for each of 81 candidates, totalling 405 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed: 6.3s
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed: 24.0s
[Parallel(n_jobs=-1)]: Done 349 tasks     | elapsed: 50.7s
[Parallel(n_jobs=-1)]: Done 405 out of 405 | elapsed: 58.8s finished
```

```
Tuned Random Forest - Accuracy: 0.9385307346326837
Tuned Random Forest - Recall: 0.6237623762376238
Tuned Random Forest - Precision: 0.9545454545454546
Tuned Random Forest - F1-score: 0.7544910179640719
```

	precision	recall	f1-score	support
False	0.94	0.99	0.96	566
True	0.95	0.62	0.75	101
accuracy			0.94	667
macro avg	0.95	0.81	0.86	667
weighted avg	0.94	0.94	0.93	667



```
In [492]: # The best parameters after tuning
best_params = rf_grid.best_params_
print("Best Hyperparameters for Random Forest Model:")
print(best_params)
```

```
Best Hyperparameters for Random Forest Model:
{'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}
```

Model Evaluation

The tuned Random Forest model shows excellent performance with a high accuracy (94%) and very good precision (95%)

Conclusion & Recommendation

The Random Forest Classifier emerged as the best model for churn prediction, achieving the highest accuracy (93.85%) and a good balance between recall (0.62) and precision (0.95) after hyperparameter tuning. The Decision Tree Classifier also performed well with an accuracy of 92%.

Models like KNN, SVM, and Naive Bayes showed poorer performance and are not recommended for this task.

Recommendation:

- Use Random Forest as the primary model due to its balanced performance and high accuracy.
- Consider further tuning.

Based on the data and analysis, one of the key attributes predictive of churn is the number of customer service calls. Customers who have had more interactions with customer service tend to be more likely to churn

Customers who make frequent customer service calls could be facing unresolved issues. By analyzing the types of issues raised, SyriaTel can prioritize improvements in the areas that matter most to customers

In []: